

Privacy and Security FAQ

Data Handling, Runtime Isolation, and University Pilot Security

Agent Failure Labs University Pilot Pack

FAQ Summary

Agent Failure Labs are designed for controlled educational use. Lab scenarios use fictitious organizations, fictitious people, and synthetic sensitive information. The platform is intended to minimize student personal data collection while preserving the trace evidence instructors need for teaching, feedback, and grading.

1. Purpose of This FAQ

This FAQ is intended for instructors, university administrators, legal reviewers, privacy officers, and IT security teams evaluating Agent Failure Labs for course or pilot use.

It answers common questions about:

- student data collection;
- trace and lab activity logging;
- handling of student personally identifiable information;
- data retention and deletion;
- third-party service use;
- runtime isolation;
- platform abuse controls;
- university pilot security considerations.

Important Note

This FAQ is a product and pilot-readiness document, not a formal legal certification. Institutions with specific legal, procurement, or compliance requirements should review those requirements with the Agent Failure team before adoption.

2. Platform Overview

2.1 What is Agent Failure?

Agent Failure is a hosted educational cyber range for AI agent security. Students complete controlled labs involving intentionally vulnerable AI agents. The labs are designed to teach concepts such as prompt injection, tool misuse, trace analysis, sensitive-data disclosure, and agent security mitigations.

2.2 Is this a real attack environment?

No. Agent Failure Labs use controlled, synthetic environments. The target systems, organizations, people, email messages, secrets, and sensitive data used in the labs are fictitious.

Students are not authorized to attack real systems, real users, production services, or third-party platforms.

2.3 What is Lab 1 about?

Lab 1, *OpsMail Assistant: Indirect Prompt Injection via Poisoned Inbox*, teaches students how attacker-controlled email content can influence an AI agent that reads inbox messages. The lab uses fictitious protected information, such as a fictional manager’s home address, to demonstrate a confidentiality failure in a safe environment.

3. Data Collection

3.1 What student data may be collected?

Depending on pilot configuration, the platform may collect or process the following categories of data:

Data Type	Purpose
Account identifier	Allows students to log in and associate progress with a course.
Email address or username	Used for account access if required by the pilot configuration.
Course or class membership	Allows instructors to manage student access and progress.
Lab progress	Records whether students started, completed, or reset a lab.
Student lab inputs	Includes prompts, attacker emails, and other inputs entered into the lab.
Assistant outputs	Includes responses generated during the lab.
Trace events	Records tool calls, tool outputs, event timeline entries, and success/failure events.
Report submissions	Stored if the platform supports report upload or report generation.
Operational logs	Used for debugging, security, reliability, and abuse prevention.

3.2 What data is not required for standard pilot use?

Standard pilot use should not require students to provide:

- government identification numbers;
- home addresses;

- phone numbers;
- payment card details;
- personal API keys;
- campus network credentials;
- local device credentials;
- real sensitive personal information.

3.3 Can students use pseudonymous identifiers?

Where institution policy permits, students may be able to use pseudonymous or course-specific identifiers rather than full legal names. This can reduce privacy risk while still allowing instructors to associate work with students for grading.

The exact identity model should be agreed during pilot setup.

4. Student PII Minimization

4.1 How does Agent Failure reduce student PII collection?

The platform is intended to collect only the information needed to:

- provide student access to labs;
- associate attempts with a course;
- enable instructor review;
- support grading and feedback;
- maintain platform reliability and security.

For pilots, the recommended configuration is to use the minimum student identifier needed for course administration.

4.2 Should students enter real personal data into labs?

No. Students should not enter real personal data, real secrets, real credentials, real addresses, real customer information, or third-party private information into the lab environment.

Lab scenarios are designed to use synthetic data. Instructors should remind students that the lab is a controlled educational environment.

Student Safety Instruction

Students should treat the lab as a simulated environment and should not enter real personal data, credentials, confidential university information, employer information, or third-party data into prompts, emails, reports, or lab inputs.

5. Lab Trace Data

5.1 What are traces?

Traces are structured records of what happened during a lab attempt. They are used for teaching, debugging, grading, and security analysis.

A trace may include:

- student prompts;
- attacker emails written inside the lab;
- assistant responses;
- tool calls;
- tool outputs;
- event timeline entries;
- success or failure events;
- timestamps;
- lab session identifiers.

5.2 Why are traces collected?

Trace collection is central to the educational purpose of Agent Failure Labs. It allows students and instructors to understand not only whether an exploit succeeded, but how it happened.

Traces help students:

- reconstruct the exploit chain;
- identify where untrusted content entered the agent context;
- support claims in their security reports;
- distinguish observation from interpretation;
- propose better mitigations.

Traces help instructors:

- verify lab completion;
- review student attempts;
- grade reports objectively;
- identify common failure modes;
- provide targeted feedback.

5.3 Who can see student traces?

In a standard course setup, traces should be visible to:

- the student who generated the trace;
- the instructor or teaching staff for the course;
- authorized platform administrators for support, reliability, and security purposes.

Student traces should not be visible to other students unless the instructor intentionally uses anonymized examples for teaching.

6. Data Retention and Deletion

6.1 How long is pilot data retained?

Retention should be agreed with the instructor or institution before the pilot.

A recommended default for early pilots is:

- lab activity and trace data retained during the course or pilot;
- data retained for a short review period after course completion;
- course data deleted or anonymized after the agreed retention period.

Suggested Default Retention Policy

For university pilots, Agent Failure may retain course data for up to 90 days after course completion unless otherwise agreed. Instructors may request earlier deletion or an institution-specific retention period.

6.2 Can course data be deleted after a pilot?

Yes. Instructors or institutional contacts should be able to request deletion of course data after the pilot, subject to operational, legal, and security requirements.

A deletion request should specify:

- course or pilot identifier;
- date range;
- whether deletion should apply to all students or selected students;
- whether anonymized aggregate usage metrics may be retained;
- institutional contact authorizing the deletion.

6.3 Can individual student data be deleted?

Where technically feasible and consistent with course policy, individual student data may be deleted upon authorized request.

For graded coursework, deletion timing should be coordinated with the instructor to avoid interfering with assessment records.

7. Use of Student Data

7.1 Is student data sold?

No. Student lab data should not be sold.

7.2 Is student lab data used to train public foundation models?

Student lab data should not be used to train public foundation models.

If third-party model APIs are used to power lab agents, those providers' data handling terms should be reviewed and disclosed as part of the pilot configuration.

7.3 Is student data used to improve Agent Failure?

Aggregated or anonymized data may be used to improve lab design, identify common failure modes, improve teaching materials, and improve platform reliability.

Examples include:

- average time to completion;
- common points where students get stuck;
- frequency of successful lab completion;
- common error categories;
- anonymized examples of failure modes.

Identifiable student work should not be used in public marketing, publications, or demonstrations without permission.

8. Third-Party Services

8.1 Does Agent Failure use third-party services?

A hosted platform may use third-party services for hosting, authentication, logging, error monitoring, analytics, and AI model inference.

The exact services should be documented for each deployment or pilot.

8.2 Do students communicate directly with LLM providers?

For university pilots, the preferred architecture is that students communicate with Agent Failure's hosted platform, and the platform communicates with any model provider server-side.

Students should not need their own API keys and should not need direct browser access to model provider APIs.

Service Category	Possible Use	Data Consideration
Cloud hosting	Runs web application and lab infrastructure	May process application and run-time data.
Authentication	User login and session management	May process identifiers such as email or username.
LLM API provider	Generates agent responses	May receive prompts, lab state, or synthetic scenario content if used server-side.
Error monitoring	Debugging and reliability	May receive error logs and limited metadata.
Analytics	Product improvement	Should be minimized and disclosed if used.
Email service	Platform notifications, if enabled	May process instructor or student email addresses.

8.3 What should be disclosed to institutions?

Before a pilot, Agent Failure should disclose:

- the main hosting provider;
- authentication provider, if any;
- model provider, if any;
- whether analytics are enabled;
- whether error monitoring is enabled;
- whether student prompts or traces are sent to any third-party service.

9. Runtime Isolation and Lab Security

9.1 Are student lab sessions isolated?

Student lab sessions should be isolated from one another. One student's actions should not alter another student's lab state, inbox, trace, or result.

The expected isolation model is:

- separate lab state per student or per attempt;
- separate synthetic inbox per lab session;
- separate trace records per session;
- restricted access to other students' data;
- resettable lab environment.

9.2 Can a student attack another student's runtime?

No. The platform should prevent students from accessing or modifying other students' lab sessions. Students are only authorized to interact with the lab interfaces assigned to them.

9.3 Can a student attack the hosting platform?

No. The rules of engagement prohibit attacks on the hosting infrastructure, backend services, authentication systems, databases, networks, or unrelated platform components. The lab is limited to the educational scenario and provided interfaces.

9.4 Does the lab access real sensitive data?

No. Lab scenarios should use fictitious organizations, fictitious people, and synthetic sensitive information. For example, Lab 1 uses a fictional manager and a fictional protected address.

9.5 Does the lab have internet access?

The runtime network configuration may vary by deployment. For educational safety, lab runtimes should be restricted to the services required for the lab.

A recommended configuration is:

- no unrestricted outbound internet access from student-controlled lab components unless required;
- no inbound access from the public internet to student sessions;
- restricted tool permissions;
- synthetic data only;
- monitoring for abuse or out-of-scope activity.

10. Instructor and Administrator Access

10.1 What can instructors see?

Depending on the platform configuration, instructors may be able to see:

- roster or class membership;
- lab start and completion status;
- success/failure indicators;
- student trace records;
- student-submitted reports;
- timestamps and progress events;
- grade export data, if enabled.

10.2 Can instructors export grades?

If grade export is enabled, instructors may be able to export completion or grading data as CSV or another supported format.

The export should include only information needed for course administration.

10.3 Can teaching assistants access student traces?

Teaching assistant access should be controlled by the instructor or course administrator. If enabled, teaching assistants may review traces and reports for grading or support.

11. Security Controls

11.1 What security controls are expected?

A hosted educational cyber range should include controls such as:

- HTTPS transport security;
- authenticated access to course areas;
- authorization checks for student and instructor data;
- per-session or per-student lab isolation;
- synthetic lab data;
- logging and monitoring;
- abuse detection or rate limiting;
- runtime reset capability;
- restricted lab tooling;
- support process for vulnerability reports.

11.2 How are lab prompts and outputs protected?

Lab prompts and outputs should be stored and accessed according to the course access model. Students should only see their own work unless otherwise configured. Instructors and authorized teaching staff may review student traces and submissions for educational and grading purposes.

11.3 What happens if a student finds a platform vulnerability?

Students should report suspected platform vulnerabilities through the designated security contact rather than attempting further exploitation.

Security Contact Placeholder

Security contact: security@agentfailure.com

12. Legal and Compliance Considerations

12.1 Is Agent Failure FERPA compliant?

Agent Failure should be designed to support privacy-conscious educational use, but formal FERPA evaluation depends on institutional use, data flows, contracts, and administrative configuration.

Do not assume formal FERPA compliance without institutional review.

12.2 Is Agent Failure GDPR compliant?

GDPR applicability depends on the location of students, institution, data processing arrangements, subprocessors, retention policy, and legal basis for processing.

If GDPR applies, the institution and Agent Failure should review required terms before deployment.

12.3 Is Agent Failure SOC 2 or ISO 27001 certified?

Unless a current certification exists, do not represent the platform as SOC 2 or ISO 27001 certified.

For early pilots, the appropriate statement is that the platform is designed with data minimization, hosted runtime isolation, and controlled educational use in mind, but formal certification status should be confirmed separately.

12.4 Does this lab involve real personal data?

No. Lab scenarios are intended to use synthetic personal data. Students should not enter real personal data.

13. Recommended University Pilot Configuration

For initial university pilots, the recommended configuration is:

Area	Recommended Configuration
Student identity	Minimal identifier or course-specific pseudonym where possible.
Lab data	Synthetic only.
Student API keys	Not required.
Runtime access	Browser-only, hosted sessions.
Trace access	Student, instructor, authorized teaching staff.
Retention	Delete or anonymize shortly after course completion.
Analytics	Disabled or minimized for pilots unless disclosed.
Public examples	Use synthetic or anonymized examples only.
Support	Instructor escalation path plus security contact.

14. Frequently Asked Questions

14.1 Do students need to install anything?

No. Students use a browser.

14.2 Do students need to provide real personal information?

Only the minimum account or course identifier required for access should be used. Students should not enter real personal data into the lab scenarios.

14.3 Are the lab targets real companies or real people?

No. Lab targets are fictitious.

14.4 Can students access other students' work?

No. Student work should be isolated and access-controlled.

14.5 Can instructors review student attempts?

Yes. Instructor review is part of the educational and grading model.

14.6 Are prompts and traces stored?

Yes, prompts, emails, assistant responses, tool calls, tool outputs, and trace events may be stored for grading, feedback, debugging, and platform reliability.

14.7 Can data be deleted after the course?

Yes, course data deletion or anonymization should be available upon authorized request, subject to agreed retention requirements.

14.8 Is student data used to train public models?

No. Student lab data should not be used to train public foundation models.

14.9 Do students need LLM provider accounts?

No. Students should not need their own LLM API keys or provider accounts.

14.10 What if a student enters real sensitive information by mistake?

The instructor or authorized institutional contact should notify the Agent Failure support or privacy contact promptly and request deletion or remediation of the affected data.

14.11 What if a student attempts out-of-scope activity?

The instructor may remove access, reset the lab, or handle the issue under course conduct policy. Serious platform abuse should be reported to Agent Failure support.

15. Data Handling Summary

Question	Summary Answer
Does the platform collect student lab activity?	Yes, for grading, trace review, and learning feedback.
Does the lab use real sensitive data?	No, lab data is synthetic.
Should students enter real personal data?	No.
Are student traces visible to classmates?	No, not by default.
Can instructors review traces?	Yes.
Is local installation required?	No.
Are students given access to real infrastructure?	No.
Can data be deleted after a pilot?	Yes, by authorized request and agreed policy.
Is data sold?	No.
Is student data used to train public models?	No.

16. Institution Review Checklist

Before adopting Agent Failure Labs, institutional reviewers may wish to confirm:

- what student identifiers are required;
- what data is collected during lab attempts;
- how long data is retained;
- how deletion requests are handled;
- which third-party services are used;
- whether LLM provider calls include student inputs or traces;
- whether analytics are enabled;
- who can access student traces;
- whether runtime isolation meets institutional expectations;
- whether the course use case requires a data processing agreement or other contract.

17. Contact Information

Contact Placeholder

General support: support@agentfailure.com
Privacy contact: privacy@agentfailure.com
Security contact: security@agentfailure.com
Platform status: <https://status.agentfailure.com>

18. End of Privacy and Security FAQ